

# Transactions, Recovery and Concurrency (I)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Air-line Reservation

- 10 available seats vs 15 travel agents.
- How do you design a robust and fair reservation system?  
  - Insufficient resources
  - Fair policy to every body
  - Robustness

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Failures

Number of factors might cause failures in user requirements processing.

1. System failure:
  - Disk failure - e.g. head crash, media fault.
  - System crash - unexpected failure requiring a reboot.
2. Program error - e.g. a divide by zero.
3. Exception conditions - e.g. no seats for your reservation.
4. Concurrency control - e.g. deadlock, expired locks.

# To handle failures correctly and efficiently

Each database user must express his requirements as a set of program units.


Each program unit is a *transaction* that either

- accesses the contents of the database, or
- changes the state of the database, from one consistent state to another

Assignment Project Exam Help  
<https://powcoder.com>  
Add WeChat powcoder

▪ Sydney → Tokyo → LA  
→ N.Y

▪ It does not make sense only partial trip has tickets



*Example transaction:* buy a ticket from Sydney to N.Y. by JAL.

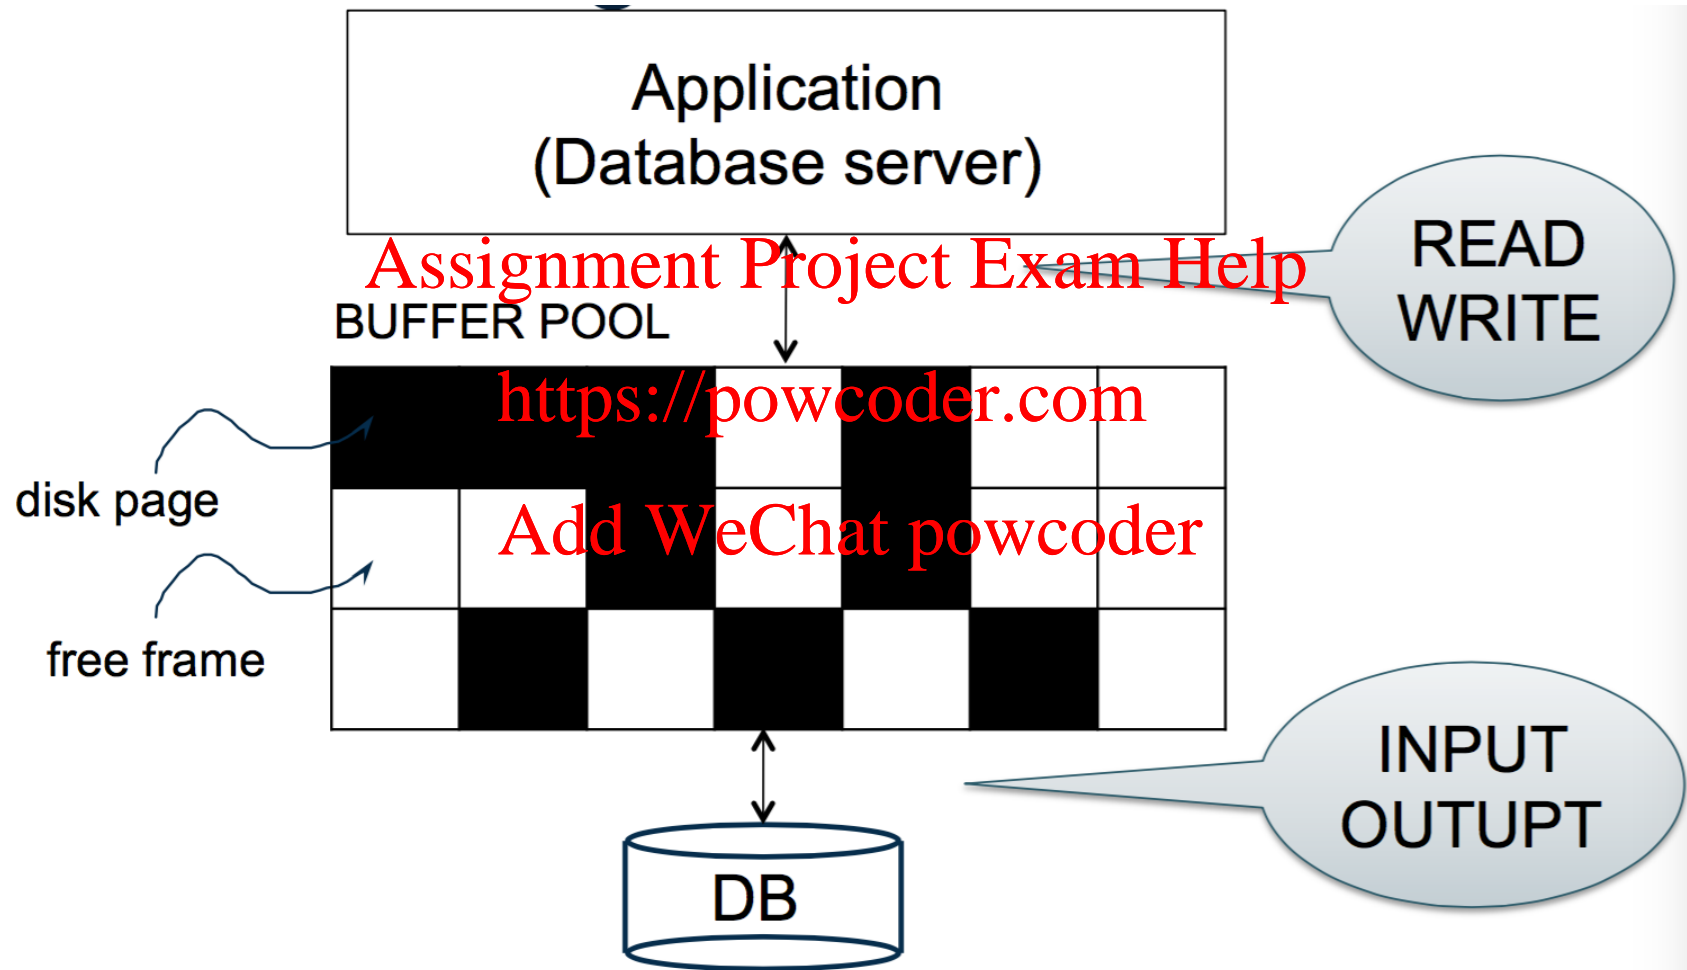
A transaction must be treated as an *atomic* unit.

# Transaction Processing

Three kinds of operations may be used in a transaction:

- *Read.*      **Assignment Project Exam Help**
- *Write.*      **<https://powcoder.com>**
- *Computation.*      **Add WeChat powcoder**

# Buffer Management in a DBMS



# Read

1. Compute the data block that contains the item to be read
2. Either **Assignment Project Exam Help**
  - find a buffer containing the block, or
  - read from disk into a buffer**<https://powcoder.com>**  
**Add WeChat powcoder**
3. Copy the value from the buffer.


# Write

1. Compute the disk block containing the item to be written,
2. Either **Assignment Project Exam Help**
  - find a buffer containing the block, or
  - read from disk into a buffer,**<https://powcoder.com>**  
**Add WeChat powcoder**
3. Copy the new value into the buffer,
4. At some point (maybe later), write the buffer back to disk.



# Processing States of a Transaction

- The typical processing states are illustrated in the figure below (E/N Fig 17.4):



- Partially committed point:*** At this point, check and enforce the correctness of the concurrent execution.
- Committed state:*** Once a transaction enters the committed state, it has concluded its execution successfully.

# Desirable Properties of Transaction Processing ACID

- Atomicity: A transaction is either performed in its entirety or not performed at all.

## Assignment Project Exam Help

- Consistency preservation: A correct execution of the transaction must take the database from one consistent state to another.  
<https://powcoder.com>
- Isolation: A transaction should not make its updates visible to other transactions until it is committed.  
Add WeChat powcoder
- Durability or permanency: Once a transaction changes the database and the changes are committed, these changes must never be lost because of subsequent failure.

# Problems without Enforcing ACID

- For a banking system,
  - If durability is not enforced, then a customer may lose a deposit.
  - If consistency preservation is not enforced, then the bank runs a high risk of bankrupt. E.g., run-over upper-Audit WeChat powcoder
- Below are the problems if atomicity and isolation are not enforced in a concurrent execution of transactions.

# Lost Update Problem (Isolation is not enforced)

- Suppose we have these two transactions,  $T_1$  and  $T_2$ :

$T_1$ :

read( $X$ )

$X \leftarrow X + N$

write( $X$ )

read( $Y$ )

$Y \leftarrow Y - N$

write( $Y$ )

$T_2$ :

read( $X$ )

$X \leftarrow X + M$

write( $X$ )

- Let us see what may happen if  $T_1$  and  $T_2$  are executed concurrently in an uncontrolled way:

Suppose initially that  $X = 100$ ;  $Y = 50$ ;  $N = 5$  and  $M = 8$ .

Database	$T_1$	$T_2$
$X = 100, Y = 50$	$X = ?, Y = ?$	$X = ?$
$X = 100, Y = 50$	read(X) $X = 100, Y = ?$	$X = ?$
$X = 100, Y = 50$	$X \leftarrow X + N$ $X = 105, Y = ?$	$X = ?$
$X = 100, Y = 50$	read(X) $X = 105, Y = ?$	$X = 100$
$X = 100, Y = 50$	$X \leftarrow X + M$ $X = 105, Y = ?$	$X = 108$
$X = 105, Y = 50$	write(X) $X = 105, Y = ?$	$X = 108$
$X = 105, Y = 50$	read(Y) $X = 105, Y = 50$	$X = 108$
$X = 108, Y = 50$	$X = 105, Y = 50$	write(X) $X = 108$
$X = 108, Y = 50$	$Y \leftarrow Y - N$ $X = 105, Y = 45$	$X = 108$
$X = 108, Y = 45$	write(Y) $X = 105, Y = 45$	$X = 108$

- At the end of  $T_1$  and  $T_2$ ,  $X$  should be 113,  $Y$  should be 45.
- The update  $X \leftarrow X + N$  has been lost

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Incorrect Summary Problem (Isolation Issue)

$T_1$	$T_3$
	$sum \leftarrow 0$
	read(A)
	$sum \leftarrow sum + A$
	$\vdots$
read(X)	
$X \leftarrow X - N$	
write(X)	
	$\vdots$
	read(X)
	$sum \leftarrow sum + X$
	read(Y)
	$sum \leftarrow sum + Y$
	$\vdots$
read(Y)	
$Y \leftarrow Y + N$	
write(Y)	
	$\vdots$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- Here the sum calculated by  $T_3$  will be wrong by  $N$ .

# The Temporary Update Problem

Database	$T_1$	$T_2$
$X = 100, Y = 50$	$X = ?, Y = ?$	$X = ?$
$X = 100, Y = 50$	read(X) $X = 100, Y = ?$	$X = ?$
$X = 100, Y = 50$	$X \leftarrow X + N$ $X = 105, Y = ?$	$X = ?$
$X = 105, Y = 50$	write(X) $X = 105, Y = ?$	$X = ?$
$X = 105, Y = 50$	<b>**FAILS**</b>	read(X) $X = 105$
$X = 105, Y = 50$		$X \leftarrow X + M$ $X = 113$

Recover from the disk

Several possibilities for what might happen next:



Database	$T_1$	$T_2$
X = 105, Y = 50		X = 113
X=100, Y=50	Case 1: DBMS undoes $T_1$	X = 113
X=113, Y=50		Write (X) X= 113

Database	$T_1$	$T_2$
X = 105, Y = 50		X = 113
X=105, Y=50	Case 2: DBMS does nothing to $T_1$	X = 113
X=113, Y=50		Write (X) X= 113

- Case 1&2, only half of  $T_1$  has been executed.
- Case 3,  $T_1$  &  $T_2$  have been lost.

Database	$T_1$	$T_2$
X = 105, Y = 50		X = 113
X= 105, Y = 50		X = 113
X=100, Y=50	Case 3: DBMS undoes $T_1$	Write (X), X= 113 X= 100

# Recover from Failures

Ensure ACID

Assignment Project Exam Help

Log-based Recovery

<https://powcoder.com>

- Undo logging

- Redo logging

- Undo/Redo logging

# System Log

- System Log

- The system needs to record the states information to recover failures correctly.

<https://powcoder.com>

- The information is maintained in a log (also called journal or audit trail).

- The system log is kept in hard disk but maintains its current contents in main memory.

# System Log

- Start transaction marker [start transaction,  $T$ ]: Records that transaction  $T$  has started execution.
- [read item,  $T, X$ ]: Records that transaction  $T$  has read the value of database item  $X$ .
- [write item,  $T, X$ , old value, new value]: Records that  $T$  has changed the value of database item  $X$  from old value to new value.
- Commit transaction marker [commit,  $T$ ]: Records that transaction  $T$  has completed successfully, and asserts that its effect can be committed (recorded permanently) to the database.
- [abort,  $T$ ]: Records that transaction  $T$  has been aborted.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# System Log (Cont'd)

- In fact some other entries (rollback, undo, redo) are also required for a recovery method.

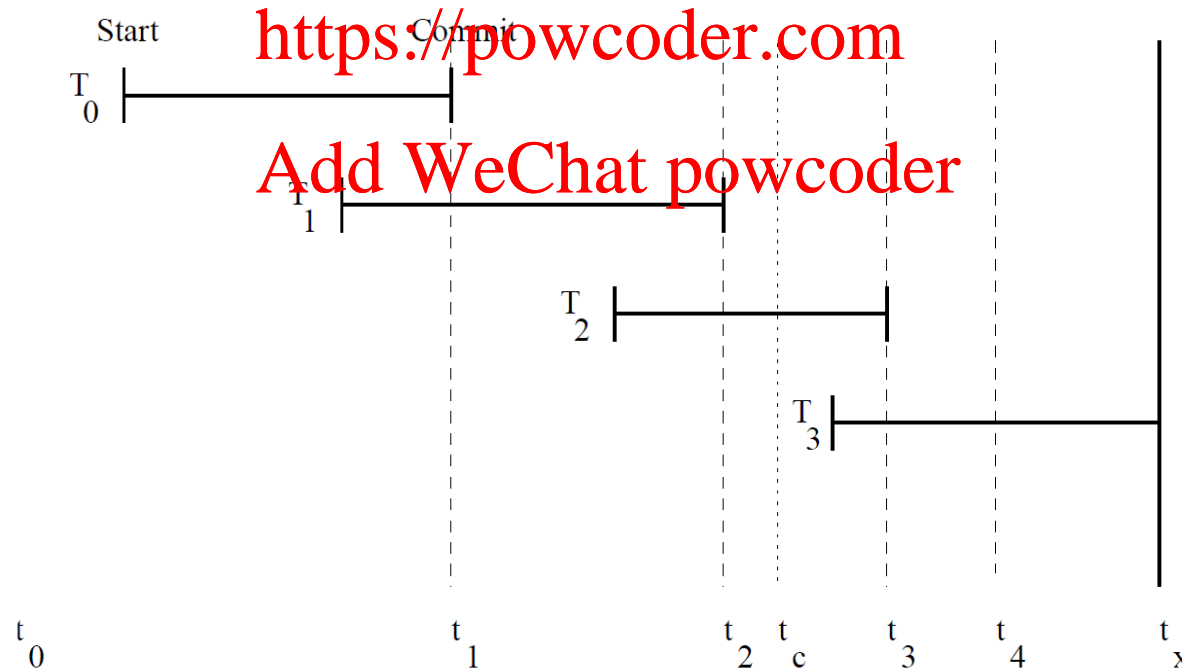
## Assignment Project Exam Help

- These entries allow the recovery manager to *rollback* an unsuccessful transaction (undo any partial updates).

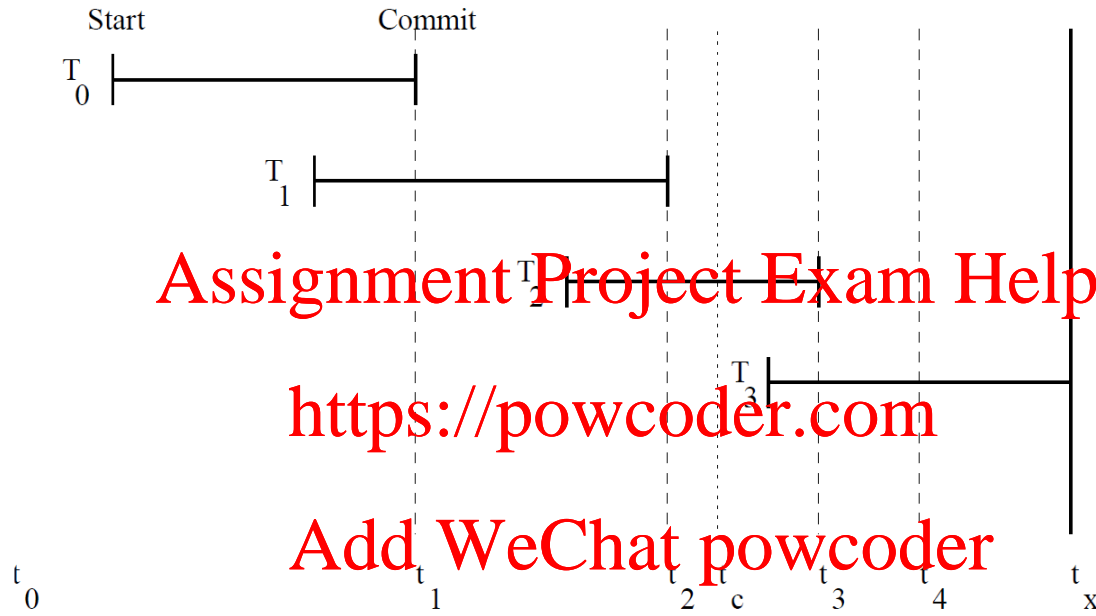
<https://powcoder.com>  
Add WeChat powcoder

# Recovery

- Let us see how the log might be used to recover from a system crash.
- The diagram below shows transactions between the last system backup and a crash.



# Recovery (Cont'd)




- The database on disk will be in a state somewhere between that at  $t_0$  and the state at  $t_x$ .
- The same is also true for log entries.

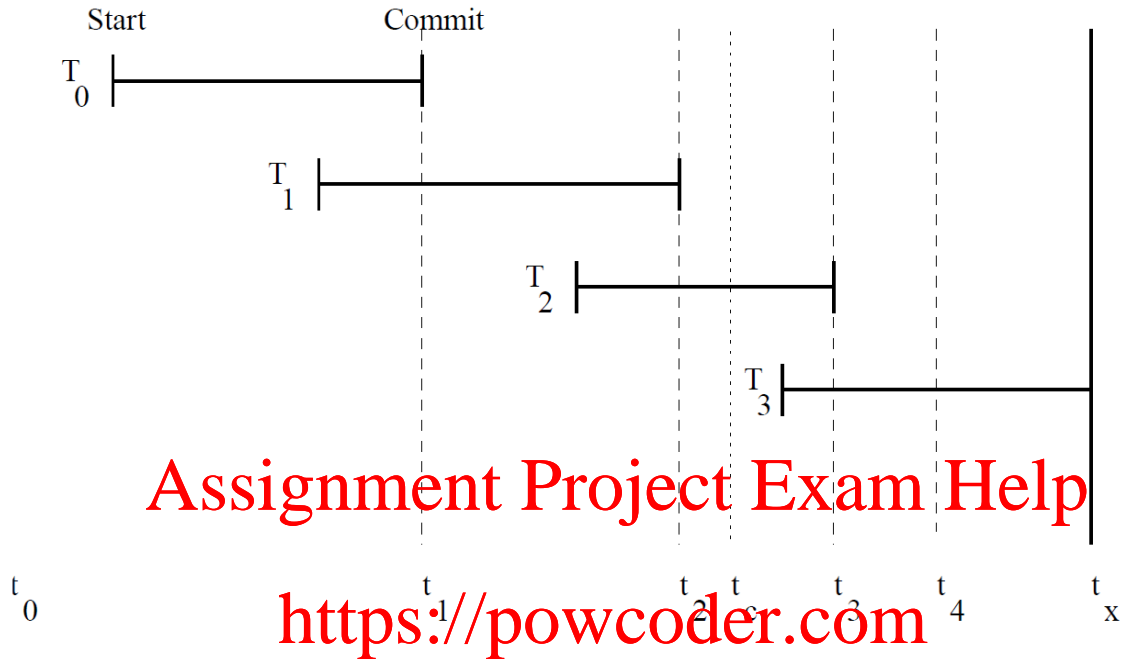
# Recovery (Cont'd)

- We will assume that the *write-ahead log strategy* is used. This means that
  - old data values must be force-written to the log (i.e. the buffer must be copied to disk) before any change can be made to the database, and
  - the transaction is regarded as committed when the new data values and the commit marker have been force-written to the log.
- Thus the log is force-written at least at  $t_1$ ,  $t_2$  and  $t_3$  in the above.





- Suppose the log was last written to disk at  $t_4$ .
- By examining the log:
  1. We know that  $T_0$ ,  $T_1$  and  $T_2$  have committed and their effects should be reflected in the database after recovery.
  2. But we do not know whether the effects of  $T_0$ ,  $T_1$  and  $T_2$  were reflected at the time of the crash.
  3. We also know that  $T_3$  has started, may have modified some data, but is not committed. Thus  $T_3$  should be undone.



- The database can be recovered by rolling back  $T_3$  using the old data values from the log, and redoing the changes made by  $T_0 \dots T_2$  using the new data values (for these committed transactions) from the log.
- Notice that instead of rolling back, the database could have been restored from the backup. This might be necessary in the event of a disk crash for example (for this reason, the log should be stored on an independent disk pack).

# Checkpoints

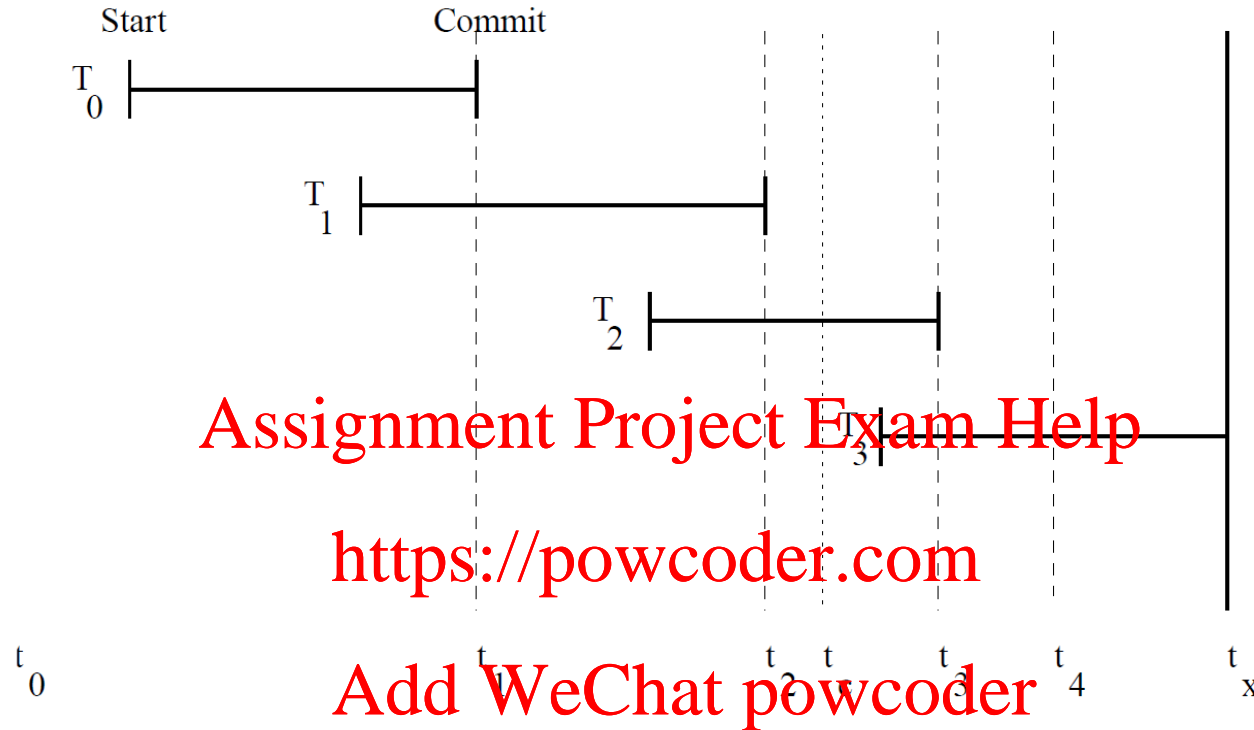
- Notice also that using this system, the longer the time between crashes, the longer recovery may take.

## Assignment Project Exam Help

- To avoid this problem, the system may take *checkpoints* at regular intervals.

<https://powcoder.com>

- To do this:
  - a *start of checkpoint* marker is written to the log, then
  - the database updates in buffers are force-written, then
  - an *end of checkpoint* marker is written to the log.



In our example, suppose a checkpoint is taken at time  $t_c$ . Then on recovery we only need redo  $T_2$ .

# Schedules of Transactions

- To fully utilise resources, desirable to interleave the operations of transactions in an appropriate way.
- For example, if one transaction is waiting for I/O to complete, another transaction can use the CPU.

<https://powcoder.com>

- A *schedule*  $S$  of the transactions  $T_1, \dots, T_n$ 
  - is a sequential ordering of the operations of  $T_1, \dots, T_n$ , and
  - preserves the ordering of operations in each transaction  $T_i$ .



(a)		(b)	
$T_1$	$T_2$	$T_1$	$T_2$
$\text{read}(X)$ $X \leftarrow X + N$ $\text{write}(X)$ $\text{read}(Y)$ $Y \leftarrow Y - N$ $\text{write}(Y)$		$\text{read}(X)$ $X \leftarrow X + N$ $\text{write}(X)$ $\text{read}(Y)$ $Y \leftarrow Y - N$ $\text{write}(Y)$	$\text{read}(X)$ $X \leftarrow X + M$ $\text{write}(X)$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

cannot swap read (X) and write (X)

(c)		(d)	
$T_1$	$T_2$	$T_1$	$T_2$
$\text{read}(X)$ $X \leftarrow X + N$  $\text{write}(X)$ $\text{read}(Y)$  $Y \leftarrow Y - N$ $\text{write}(Y)$	$\text{read}(X)$ $X \leftarrow X + M$  $\text{write}(X)$	$\text{read}(X)$ $X \leftarrow X + N$ $\text{write}(X)$  $\text{read}(Y)$ $Y \leftarrow Y - N$ $\text{write}(Y)$	$\text{read}(X)$ $X \leftarrow X + M$ $\text{write}(X)$
 <p>Incorrect</p>		 <p>correct</p>	

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- As we have seen, if operations are interleaved arbitrarily, incorrect results may occur.
- However, it is reasonable to assume that schedules (a) and (b) in the figure will give correct results (as long as the transactions are independent).
- (a) and (b) are called *serial* schedules, and we will assume that *any serial schedule is correct*.
- Notice that schedule (d) always produces the same result as schedules (a) and (b), so it should also give correct results.
- A schedule is *serializable* if it always produces the same result as some serial schedule. (see E/N 17.5.1 for a formal definition).
- Notice that schedule (c) is not serializable.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



# Scheduling Transactions

- Schedule and Complete Schedule?
- Serial schedule: Schedule that does not interleave the actions of different transactions.
- Equivalent schedules: For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second schedule.
- Serializable schedule: A schedule over a set S of transactions is equivalent to some serial execution of the set of committed transactions in S.  
(Note: If each transaction preserves consistency, every serializable schedule preserves consistency. )

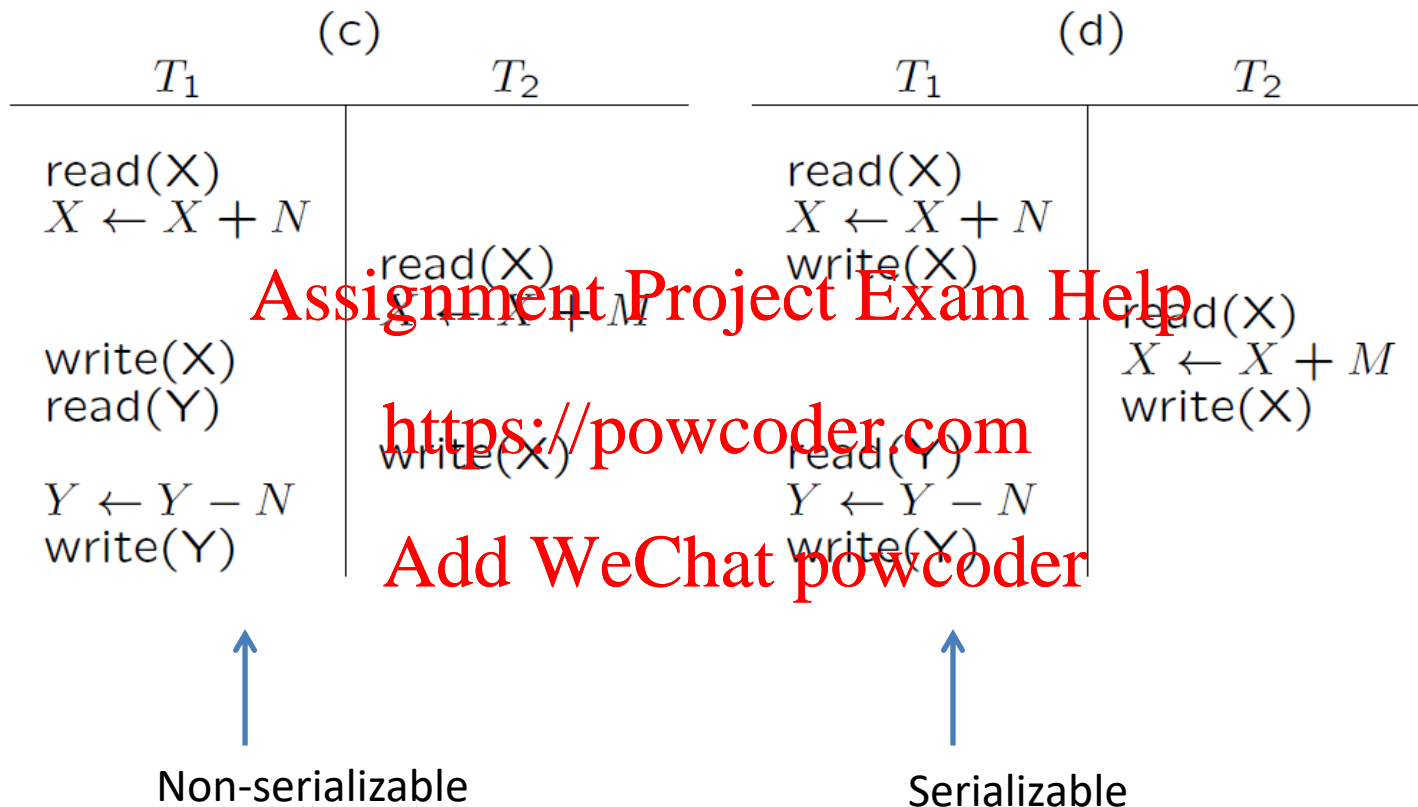
(a)		(b)	
$T_1$	$T_2$	$T_1$	$T_2$
read(X) $X \leftarrow X + N$ write(X) read(Y) $Y \leftarrow Y - N$ write(Y)		read(X) $X \leftarrow X + N$ write(X) read(Y) $Y \leftarrow Y - N$ write(Y)	read(X) $X \leftarrow X + M$ write(X)
	read(X) $X \leftarrow X + M$ write(X)	read(X) $X \leftarrow X + N$ write(X) read(Y) $Y \leftarrow Y - N$ write(Y)	

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Serial Schedule



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Scheduling Transactions (Cont.)

- Recoverable schedule (RS): Transactions commit only after (and if) all transactions whose changes they read commit.

EX1: T1.R(X), T1.W(X), T2. R(X), T2.W(X), COMMIT.T2.

EX1 is not recoverable.

EX2: T1.R(X), T1.W(X), T2. R(X), T2.W(X), COMMIT.T1. Recoverable!

- Avoid cascading aborts (ACA): Transactions read only the changes of committed transactions.

EX3: T1.R(X), T1.W(X), T2.R(X), T2.W(X), COMMIT.T1. EX3 is not ACA.

EX4: T1.R(X), T1.W(X), COMMIT.T1, T2. R(X), T2.W(X)... ACA!

- Strict schedules (SS): A value written by a transaction is not read or overwritten by other transactions until T either aborts or commits.

EX5: T1.R(X), T1.W(X), T2.W(X)... EX5 is RS and ACA but not SS.

EX6: T1.R(X), T1.W(X), COMMIT.T1, T2.W(X)... EX6 is SS.

Note: SS is ACA and ACA is RS but not vice versa.

# Check Serializability

- When there are only two transactions, there are only two serial schedules - for  $n$  transactions there will be  $n!$ .
- Fortunately there is an efficient algorithm to check whether a schedule is serializable without checking all these possibilities.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Conflict Serializable Schedules

- Two schedules are *conflict equivalent* if:
  - Involve the same actions of the same transactions
  - Every pair of conflicting actions is ordered the same way
- Schedule S is *conflict serializable* if S is conflict equivalent to some serial schedule

Assignment Project Exam Help

https://powcoder.com

Add WeChat powcoder

(d)		(a)	
$T_1$	$T_2$	$T_1$	$T_2$
read(X) $X \leftarrow X + N$ write(X)	read(X) $X \leftarrow X + M$ write(X)	read(X) $X \leftarrow X + N$ write(X) read(Y) $Y \leftarrow Y - N$ write(Y)	read(X) $X \leftarrow X + M$ write(X)
read(Y) $Y \leftarrow Y - N$ write(Y)			

# View Serializability

- Schedules S1 and S2 are *view equivalent* if:
  - If  $T_i$  reads initial value of A in S1, then  $T_i$  also reads initial value of A in S2
  - If  $T_i$  reads value of A written by  $T_j$  in S1, then  $T_i$  also reads value of A written by  $T_j$  in S2
  - If  $T_i$  writes final value of A in S1, then  $T_i$  also writes final value of A in S2
- A schedule is *view serializable* if view equivalent to a serial schedule.

T1:	R(A)	W(A)
T2:	W(A)	
T3:		W(A)

T1:	R(A),W(A)	
T2:		W(A)
T3:		W(A)

# Properties of Serizability

- View Serializability does not have monotonic property; that is, a schedule is view serializable but its sub-schedule may not necessarily view serializable.

Example:

Assignment Project Exam Help



- If no blind writes, conflict serializability is equivalent to view serializability .



# Check Conflict Serializability

- *Algorithm*

Step 1: Construct a *schedule* (or *precedence*) graph – a *directed graph*.

Assignment Project Exam Help

Step 2: Check if the graph is cycle.


<https://powcoder.com>

- Cyclic: non-serializable.
- Acyclic: serializable.

Add WeChat powcoder

- A *directed graph*  $G = (V, A)$  consists of
  - a vertex set  $V$ , and
  - an arc set  $A$  such that each arc connects two vertices.
- $G$  is *cyclic* if it contains a directed cycle.

<https://powcoder.com>  
Add WeChat powcoder



Cyclic Graph

# Construct a Schedule Graph $G_S = (V, A)$ for a schedule $S$

1. A vertex in  $V$  represents a transaction.
2. For two vertices  $T_i$  and  $T_j$ , an arc  $T_i \rightarrow T_j$  is added to  $A$  if
  - there are two *conflicting* operations  $O_1 \in T_i$  and  $O_2 \in T_j$ ,
  - in  $S$ ,  $O_1$  is before  $O_2$ .

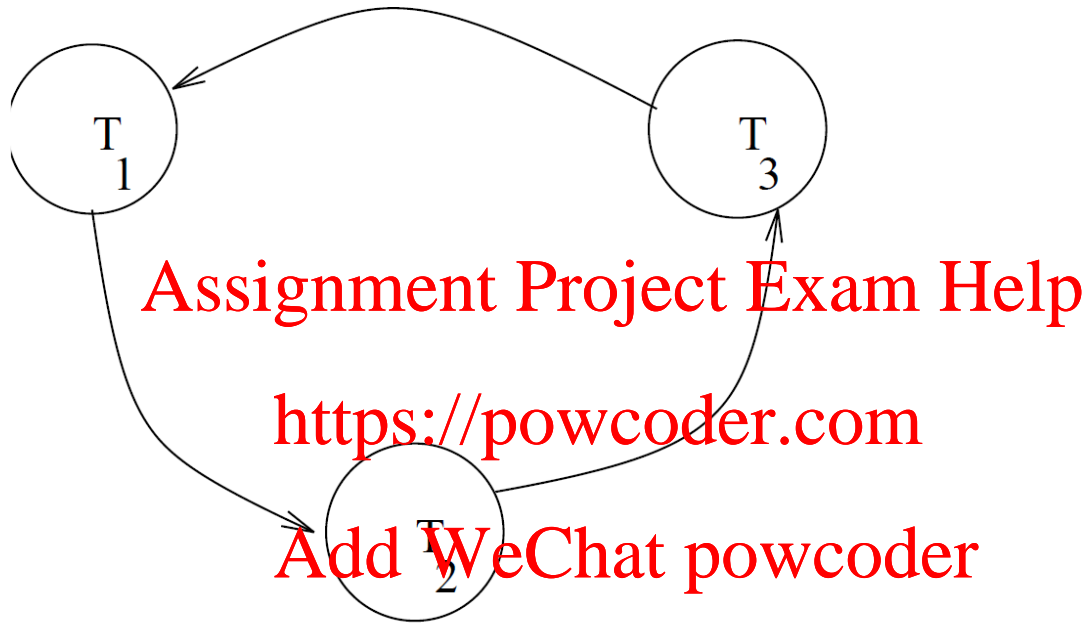
**Add WeChat powcoder**

Two operations  $O_1$  and  $O_2$  are *conflicting* if

- they are in different transactions but on the same data item,
- one of them must be a write.

### Example 1:

Schedule	$T_1$	$T_2$	$T_3$
read(A)	read(A)		
read(B)		read(B)	
$A \leftarrow f_1(A)$	$A \leftarrow f_1(A)$		
read(C)			read(C)
$B \leftarrow f_2(B)$	$B \leftarrow f_2(B)$	$B \leftarrow f_2(B)$	
write(B)		write(B)	
$C \leftarrow f_3(C)$			$C \leftarrow f_3(C)$
write(C)			write(C)
write(A)	write(A)		
read(B)			read(B)
read(A)		read(A)	
$A \leftarrow f_4(A)$		$A \leftarrow f_4(A)$	
read(C)	read(C)		
write(A)		write(A)	
$C \leftarrow f_5(C)$	$C \leftarrow f_5(C)$		
write(C)	write(C)		
$B \leftarrow f_6(B)$			$B \leftarrow f_6(B)$
write(B)			write(B)




## Example 2:

Schedule	$T_1$	$T_2$	$T_3$
read(A)	read(A)		
$A \leftarrow f_1(A)$	$A \leftarrow f_1(A)$		
read(C)	read(C)		
write(A)	write(A)		
$A \leftarrow f_2(A)$	$A \leftarrow f_2(A)$		
read(B)		read(B)	
write(C)	write(C)		
read(A)		read(A)	
read(C)			read(C)
$B \leftarrow f_3(B)$		$B \leftarrow f_3(B)$	
write(B)		write(B)	
$C \leftarrow f_4(C)$			$C \leftarrow f_4(C)$
read(B)			read(B)
write(C)			write(C)
$A \leftarrow f_5(A)$		$A \leftarrow f_5(A)$	
write(A)		write(A)	
$B \leftarrow f_6(B)$			$B \leftarrow f_6(B)$
write(B)			write(B)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- Unfortunately, testing for serializability on the fly is not practical.
- Instead, a number of protocols have been developed which ensure that if every transaction obeys the rules, then *every* schedule will be serializable, and thus correct.

- SS is serializable?

➤ irrelevant!

Example:

T1.R(X), T2.R(X), T1.W(X), COMMIT.T1,  
T2.W(X), COMMIT.T2

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder