

# Transactions, Recovery and Concurrency (II)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder  
Concurrency Control

# Concurrency Control Methods

- **Locking Mechanism**

The idea of locking some data item  $X$  is to:

- give a transaction exclusive use of the data item  $X$ ,
- do not restrict the access of other data items.

**Add WeChat powcoder**

This prevents one transaction from changing a data item currently being used in another transaction.

- We will discuss a simple locking scheme which locks individual items, using read and write locks

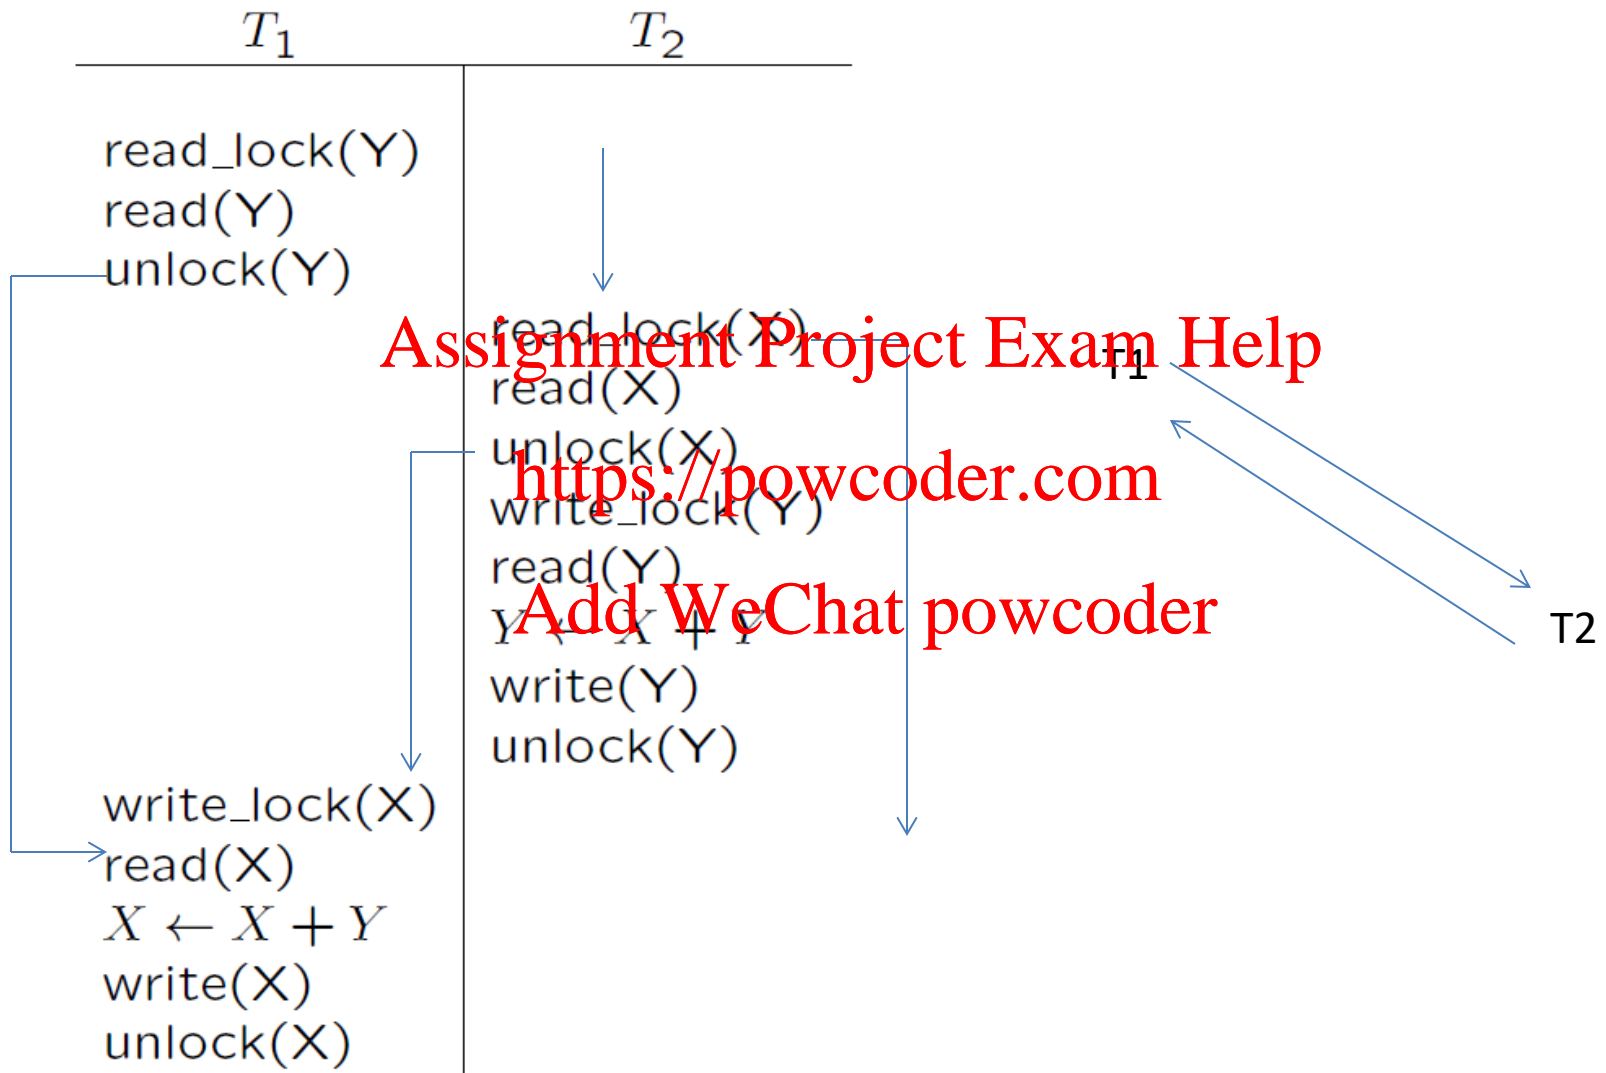
# Locking Rules

- In this schema, every transaction  $T$  must obey the following rules.
- 1) If  $T$  has only one operation (read/write) manipulating an item  $X$ :
  - obtain a read lock on  $X$  before reading it,
  - obtain a write lock on  $X$  before writing it,
  - unlock  $X$  when done with it.
- 2) If  $T$  has several operations manipulating  $X$ :
  - obtain one proper lock only on  $X$ :  
a read lock if all operations on  $X$  are reads;  
a write lock if one of these operations on  $X$  is a write.
  - unlock  $X$  after the last operation on  $X$  in  $T$  has been executed.

# Locking Rules (cont.)

- In this scheme,
  - Several read locks can be issued on the same data item at the same time.
  - A read lock and a write lock cannot be issued on the same data item at the same time, neither two write locks.
- This still does not guarantee serializability.

Example: Based on E/N Fig 18.3.



# Two Phase Locking (2PL)

- To guarantee serializability, transactions must also obey the *two-phase locking protocol*:

## Assignment Project Exam Help

- Growing Phase: all locks for a transaction must be obtained before any locks are released, and
- Shrinking Phase: gradually release all locks (once a lock is released no new locks may be requested).

<https://powcoder.com>

Add WeChat powcoder

# Two Phase Locking (2PL) (Cont.)

*Example:* Based on E/N Fig 18.4.

$T_1$

---

read\_lock(Y)  
read(Y)  
write\_lock(X)  
unlock(Y)  
read(X)  
 $X \leftarrow X + Y$   
write(X)  
unlock(X)

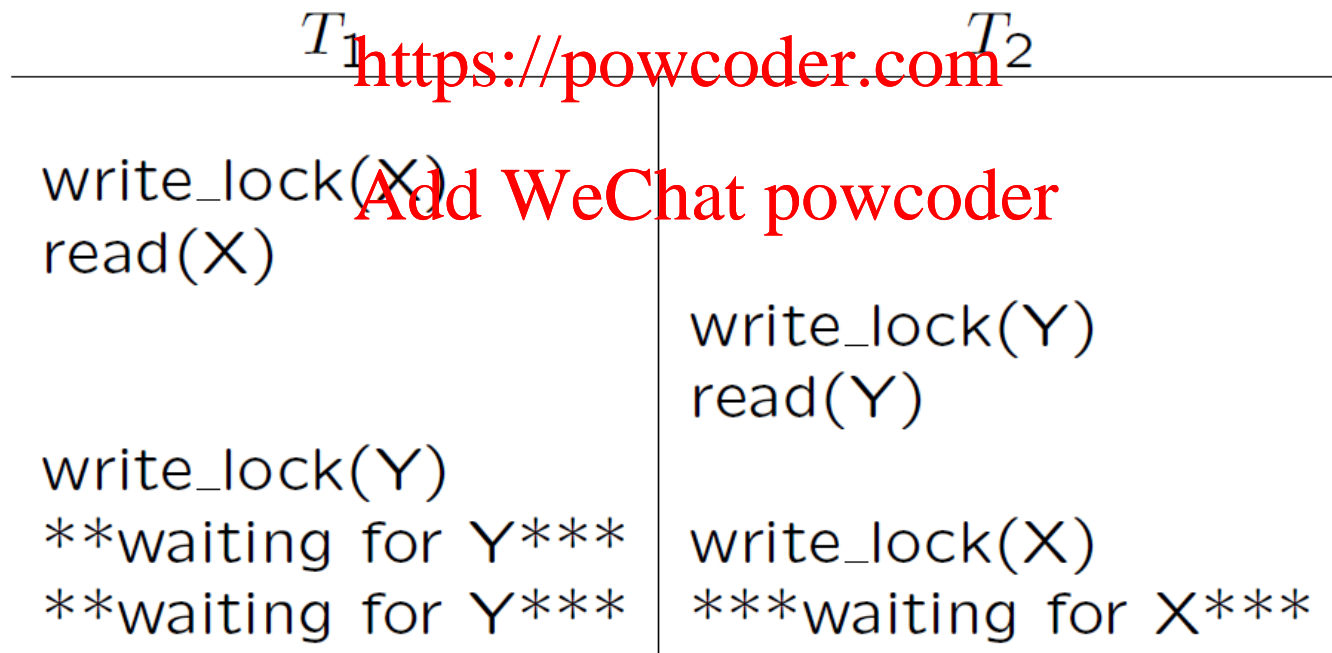
Assignment Project Exam Help  
<https://powcoder.com>  
Add WeChat powcoder

- Locking thus provides a solution to the problem of correctness of schedules.

Two phase locking ensures conflict serializability

# Deadlock

- A problem that arises with locking is **deadlock**.
- Deadlock occurs when two transactions are each waiting for a lock on an item held by the other.

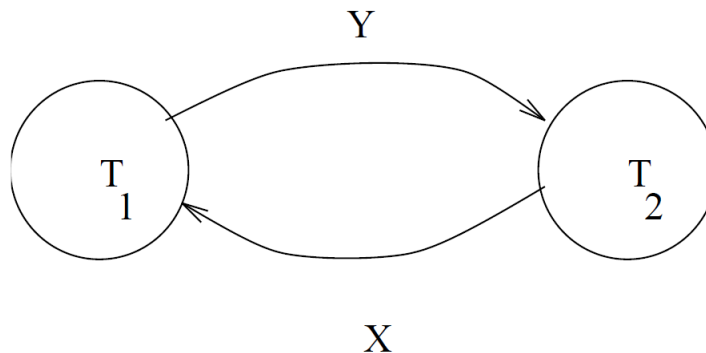




# Deadlock Check

- Create the *wait-for graph* for currently active transactions:
  - create a vertex for each transaction; and
  - an arc from  $T_i$  to  $T_j$  if  $T_i$  is waiting for an item locked by  $T_j$ .
- If the graph has a cycle, then a deadlock has occurred.

*Example.*



# Several methods to deal with deadlocks

- deadlock detection
  - periodically check for deadlocks, abort and rollback some transactions (restart them later). This is a good choice if transactions are very short or very independent.

<https://powcoder.com>

Add WeChat powcoder

# Several methods to deal with deadlocks (Cont.)

- deadlock prevention - Assign priorities based on timestamps.  
Assume  $T_i$  wants a lock that  $T_j$  holds. Two policies are possible:

## Assignment Project Exam Help

- Wait-Die: If  $T_i$  has higher priority,  $T_i$  waits for  $T_j$ ; otherwise  $T_i$  aborts

<https://powcoder.com>

- Wound-wait: If  $T_i$  has higher priority,  $T_j$  aborts; otherwise  $T_i$  waits

Add WeChat powcoder

- If a transaction re-starts, make sure it has its original timestamp

# Timestamp ordering

- The idea here is:
  - to assign each transaction a timestamp (e.g. start time of transaction), and
  - to ensure that the schedule used is equivalent to executing the transactions in timestamp order.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- Each data item,  $X$ , is assigned
  - a read timestamp,  $read\ TS(X)$  - the latest timestamp of a transaction that read  $X$ , and
  - a write timestamp,  $write\ TS(X)$  - the latest timestamp of a transaction that write  $X$ .

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- These are used in read and write operations as follows. Suppose the transaction timestamp is  $T$ .

*read(X)* **Assignment Project Exam Help**

```
If T >= write_TS(X) then
  { execute read(X);
    if T >= read_TS(X) then
      read_TS(X) <- T }
else
  rollback the transaction and restart
```

*write(X):*

```
If T >= read_TS(X) and T >= write_TS(X) then
  { execute write(X); write_TS(X) <- T }
else
  rollback and restart
```

- **Thomas' write rule:**

*write(X)*: Assignment Project Exam Help

```
If T < read_TS(X) then
    rollback and restart
else if T < write_TS(X) then
    ignore the write
else
    { execute write(X);
      write_TS(X) <- T }
```

$T_1$   
read (x)

$T_2$   
read (y)  
write (y)

read (z)

Write (z)

$T_3$   
read (z)  
write (z)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

$r\_TS(x) = 0 \rightarrow 1$

$w\_TS(x) = 0$

$r\_TS(y) = 0 \rightarrow 2$

$w\_TS(y) = 0 \rightarrow 2$

$r\_TS(z) = 0 \rightarrow 1 \rightarrow 3$

$w\_TS(z) = 0 \rightarrow 3$

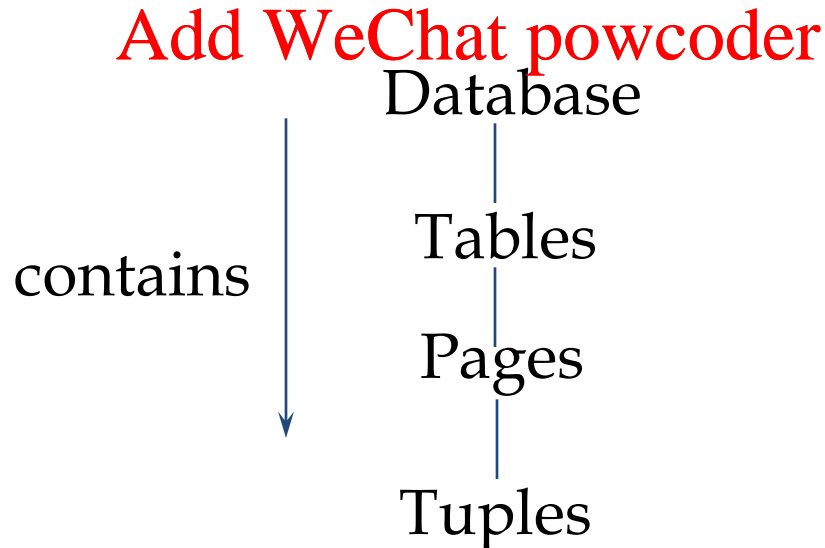


- Some problems:

- Cyclic restart: There is no deadlock, but a kind of livelock can occur - some transactions may be constantly aborted and restarted.
- Cascading rollback: When a transaction is rolled back, so are any transactions which read a value written by it, and any transactions which read a value written by them, etc. This can be avoided by not allowing transactions to read values written by uncommitted transactions (make them wait).

# Multiple-Granularity Locks

- Hard to decide what granularity to lock (tuples vs. pages vs. tables).
- Shouldn't have to decide!
- Data “containers” are nested:



# Solution: New Lock Modes, Protocol

- Allow Xacts to lock at each level, but with a special protocol using new “intention” locks:

## Assignment Project Exam Help

- ❖ Before locking an item, Xact must set “intention locks” on all its ancestors.

- ❖ For unlock, go from specific to general (i.e., bottom-up).

- ❖ **SIX mode**: Like S & IX at the same time.

	--	IS	IX	S	X
--	✓	✓	✓	✓	✓
IS	✓	✓	✓	✓	
IX	✓	✓	✓		
S	✓	✓		✓	
X	✓				

# Multiple Granularity Lock Protocol

- Each Xact starts from the root of the hierarchy.
- To get S or X lock on a node, must hold IS or IX on parent node.  
<https://powcoder.com>  
**Assignment Project Exam Help**
- What if Xact holds SIX on parent? S on parent?
- To get X or IX or SIX on a node, must hold IX or SIX on parent node.  
**Add WeChat powcoder**
- Must release locks in bottom-up order.

Protocol is correct in that it is equivalent to directly setting locks at the leaf levels of the hierarchy.

# Examples

- T1 scans R, and updates a few tuples:
  - T1 gets an SIX lock on R, then repeatedly gets an S lock on tuples of R, and occasionally upgrades to X on the tuples.
- T2 uses an index to read only part of R:
  - T2 gets an IS lock on R, and repeatedly gets an S lock on tuples of R.
- T3 reads all of R:
  - T3 gets an S lock on R.
  - OR, T3 could behave like T2; can use **lock escalation** to decide which.

	--	IS	IX	S	X
--	✓	✓	✓	✓	✓
IS	✓	✓	✓	✓	
IX	✓	✓	✓		
S	✓	✓		✓	
X	✓				

# Dynamic Databases

Sailors (sid: integer, *sname*: string, *rating*: integer, *age*: real)

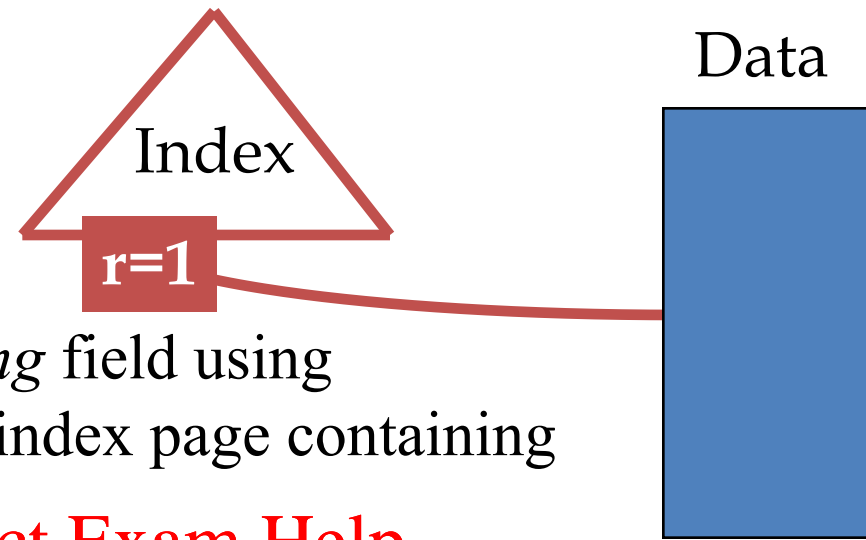
Reserves (sid: integer, bid: integer, day: dates, *rname*: string)

- If we relax the assumption that the DB is a fixed collection of objects, even Strict 2PL will not assure serializability:
  - T1 locks all pages containing sailor records with *rating* = 1, and finds oldest sailor (say, *age* = 71).
  - Next, T2 inserts a new sailor; *rating* = 1, *age* = 96.
  - T2 also deletes oldest sailor with *rating* = 2 (and, say, *age* = 80), and commits.
  - T1 now locks all pages containing sailor records with *rating* = 2, and finds oldest (say, *age* = 63).
- No consistent DB state; however T1 “correctly” gets through!

# The Problem

- T1 implicitly assumes that it has locked the set of all sailor records with *rating* = 1.
  - Assumption only holds if no sailor records are added while T1 is executing!
  - Need some mechanism to enforce this assumption. (Index locking and predicate locking.)
- Example shows that conflict serializability guarantees serializability only if the set of objects is fixed!

# Index Locking



- If there is a dense index on the *rating* field using Alternative (2), T1 should lock the index page containing the data entries with *rating* = 1.
  - If there are no records with *rating* = 1, T1 must lock the index page where such a data entry *would* be, if it existed!
- If there is no suitable index, T1 must lock all pages, and lock the file/table to prevent new pages from being added, to ensure that no new records with *rating* = 1 are added.



# Predicate Locking

- Grant lock on all records that satisfy some logical predicate, e.g. *age > 2\*salary*.
- Index locking is a special case of predicate locking for which an index supports efficient implementation of the predicate lock.
- What is the predicate in the sailor example?
- In general, predicate locking has a lot of locking overhead.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Locking in B+ Trees

- How can we efficiently lock a particular leaf node?
  - Btw, don't confuse this with multiple granularity locking!
- One solution: Ignore the tree structure, just lock pages while traversing the tree, following 2PL.
  - Root node (and many higher level nodes) become bottlenecks because every tree access begins at the root.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

# Two Useful Observations

- Higher levels of the tree only direct searches for leaf pages.
- For inserts, a node on a path from root to modified leaf must be locked (in X mode, of course), only if a split can propagate up to it from the modified leaf. (Similar point holds w.r.t. deletes.)
- We can exploit these observations to design efficient locking protocols that guarantee serializability even though they violate 2PL.

# A Simple Tree Locking Algorithm

- **Search:** Start at root and go down; repeatedly, S lock child then unlock parent.
- **Insert/Delete:** Start at root and go down, obtaining X locks as needed. Once child is locked, check if it is safe:
  - If child is safe, release all locks on ancestors.
- **Safe node:** Node such that changes will not propagate up beyond this node.
  - Inserts: Node is not full.
  - Deletes: Node is not half-empty.

ROOT



A

Do:

- 1) Search 38\*
- 2) Delete 38\*
- 3) Insert 45\*
- 4) Insert 25\*

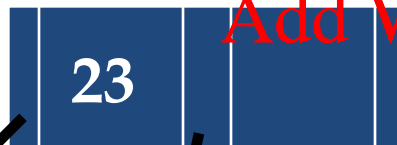


B

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



F



C

G

H

I

D

E



# A Better Tree Locking Algorithm (See Bayer-Schkolnick paper)

- **Search:** As before.
- **Insert/Delete:**
  - Set locks as if for search, get to leaf, and set X lock on leaf.
  - If leaf is not safe, release all locks, and restart Xact using previous Insert/Delete protocol.
- Gambles that only leaf node will be modified; if not, S locks set on the first pass to leaf are wasteful. In practice, better than previous alg.

ROOT

En

20

A

35

B

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

23

F

38

44

C

G

H

I

D

E

20\*

22\*

23\*

24\*

35\*

36\*

38\*

41\*

44\*

Do:

1) Delete 38\*

2) Insert 25\*

4) Insert 45\*

5) Insert 45\*,  
then 46\*

# Even Better Algorithm

- **Search:** As before.
- **Insert/Delete:**
  - Use original Insert/Delete protocol, but set IX locks instead of X locks at all nodes.
  - Once leaf is locked, convert all IX locks to X locks **top-down**: i.e., starting from node nearest to root. (Top-down reduces chances of deadlock.)

(Contrast use of IX locks here with their use in multiple-granularity locking.)





# Hybrid Algorithm

- The likelihood that we really need an X lock decreases as we move up the tree.
- Hybrid approach:


Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder



Set S locks



Set SIX locks



Set X locks

# Multiversioning

- Similar to the timestamp ordering approach; but is allowed to access “old” versions of a table.

## Assignment Project Exam Help

- A history of the values and timestamps (versions) of each item is kept.

<https://powcoder.com>

## Add WeChat powcoder

- When the value of an item is needed, the system chooses a **proper** version of the item that maintains serializability.
- This results in fewer aborted transactions at the cost of greater complexity to maintain more versions of each item.

- We will look at a scheme, several versions  $X_1, \dots, X_k$  of each data item are kept. For each  $X_i$  we also keep

## Assignment Project Exam Help

- *read*  $TS(X_i)$  - as for timestamp ordering.

<https://powcoder.com>

- *write*  $TS(X_i)$  - as for timestamp ordering.

Add WeChat powcoder

- Read and write are done as follows for a transaction  $P$  with timestamp  $T$ .

## Assignment Project Exam Help

*read(X)* <https://powcoder.com>

Find  $X_i$  s.t. write\_TS( $X_i$ ) is the  
highest write timestamp but  $\leq T$   
update read\_TS( $X_i$ ) (and do read( $X_i$ ))  
return  $X_i$  as the value for  $X$

*write(X):*

```
Find Xi s.t. write_TS(Xi) is the
    highest write timestamp but <= T
if T < read_TS(Xi) then
    rollback and restart
else
    { create a new version X(k+1) of X;
      set read_TS(X(k+1)) to T;
      set write_TS(X(k+1)) to T }
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- *Note:* Cascading rollback and cyclic restart problems can still occur, but should be reduced.
- However, there is an increased overhead in maintaining multiple versions of items.

# Optimistic scheduling

- In two-phase locking, timestamp ordering, and multiversioning concurrency control techniques, a certain degree of checking is done **before** a database operation can be executed.

Assignment Project Exam Help

- The idea here is to push on and hope for the best!

<https://powcoder.com>

- No checking is done while the transaction is executing.

Add WeChat powcoder

- The protocol has three phases.
  - read phase - A transaction can read data items from the database into local variables. However, updates are applied only to local copies of the data items kept in the transaction workspace.
  - validation phase - checks are made to ensure that serializability is not violated,
  - write phase -if validation succeeds, updates are applied and the transaction is committed. Otherwise, the updates are discarded and the transaction is restarted.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- A scheme uses timestamps and keeps each transaction's
  - read-set - the set of items read by the transaction,
  - write-set - the set of items written by the transaction.

Assignment Project Exam Help

<https://powcoder.com>

- During validation, we check that the transaction does not interfere with any transaction that is committed or currently validating.



- Each transaction  $T$  is assigned 3 timestamps:  
 $Start(T)$ ,  $Validation(T)$ ,  $Finish(T)$ .
- To pass the validation test for  $T$ , one of the following must be true:
  - 1.  $Finish(S) < Start(T)$ ; or
  - 2. for  $S$  s.t.  $Start(T) < Finish(S)$ , then
    - a) write set of  $S$  is disjoint from the read set of  $T$ , and
    - b)  $Finish(S) < Validation(T)$ .

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- Optimistic control is a good option if there is not much interaction between transactions.
- **Note:** Our earlier treatment of recovery methods largely ignored concurrency issues.

<https://powcoder.com>

Add WeChat powcoder

# 2PL vs. TSO vs. MV vs. OP

- A Comparison among two-phase locking (2PL), timestamp ordering (TSO), multiversioning (MV), optimistic (OP) concurrency control techniques.

**Assignment Project Exam Help**

- MV should provide the greatest concurrency degree (in average). However, we need to maintain multiversions for each data item.

**Add WeChat powcoder**

- 2PL can offer the second greatest concurrency degree (in average); but will result in deadlocks. To resolve the deadlocks, either
  - need additional computation to detect deadlocks and to resolve the deadlocks, or
  - reduce the concurrency degree to prevent deadlocks by adding other restrictions.

# 2PL vs. TSO vs. MV vs. OP (cont.)

- If most transactions are very short, we can use 2PL + deadlock detection and resolution.

## Assignment Project Exam Help

- TSO has a less concurrency degree than that of 2PL if a proper deadlock resolution is found. However, TSO does not cause deadlocks. Other problems, such as cyclic restart and cascading rollback, will appear in TSO.
- If there are not much interaction between transactions, OP is a very good choice. Otherwise, OP is a bad choice.